





Neural Networks

NN – Intro

- ML method used for regression, classification, etc.
- Inspired by the human brain.
- Very flexible.
- Not easily interpretable.
- Very famous nowadays because of its good performance on:
 - Images.
 - Text.
 - Audio.
 - ...




1 Neuron




1 Neuron

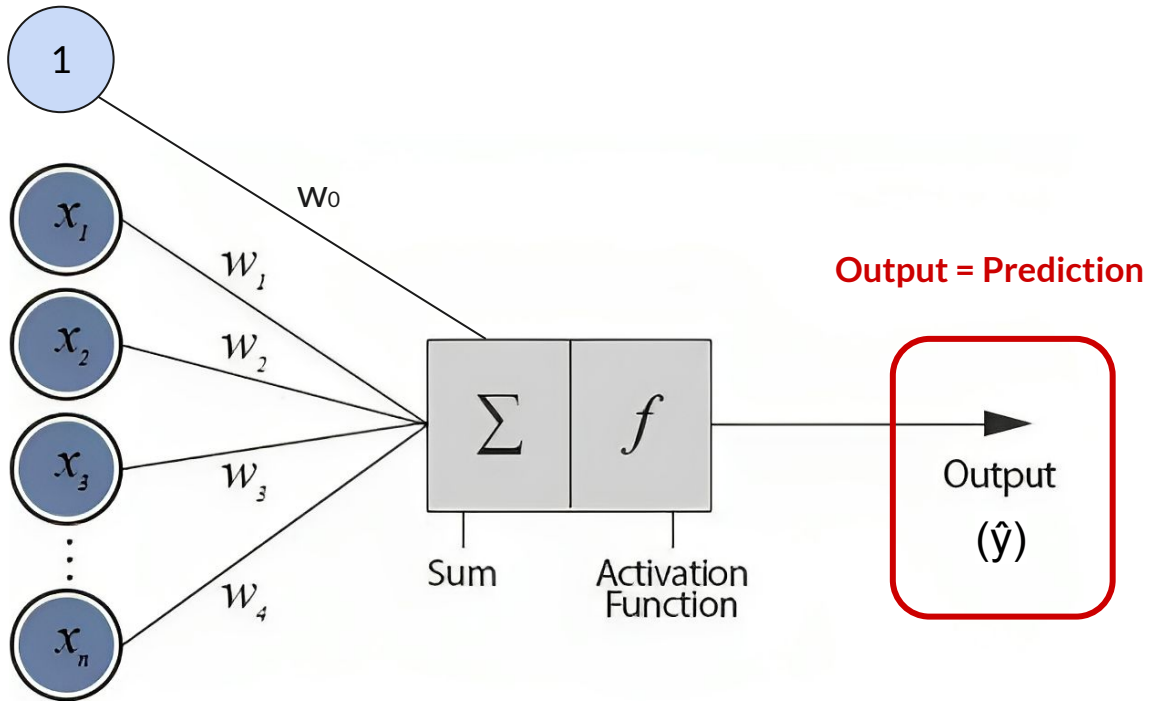
Inputs = Features



1 Neuron




1 Neuron



1 Neuron

Forward propagation:


- For each input row:
 - Multiply features with weights and sum.
 - Prediction = Activation function of previous sum.
- Calculate loss function (prediction error).




1 Neuron

Backward propagation:

- Update each weight for optimizing loss function.
- i.e., apply gradient descent.



1 Neuron








NN – Activation Function

- Aren't we just doing the same as with LR? Yes, kind of.
- Activation functions are used for learning non-linearities between features and target.
- By applying a non-linear function to a linear combination of weights and features.

NN – Common Activation Functions

- Linear
- Sigmoid
- Tanh
- ReLU
- Leaky ReLU
- ELU
- Softmax



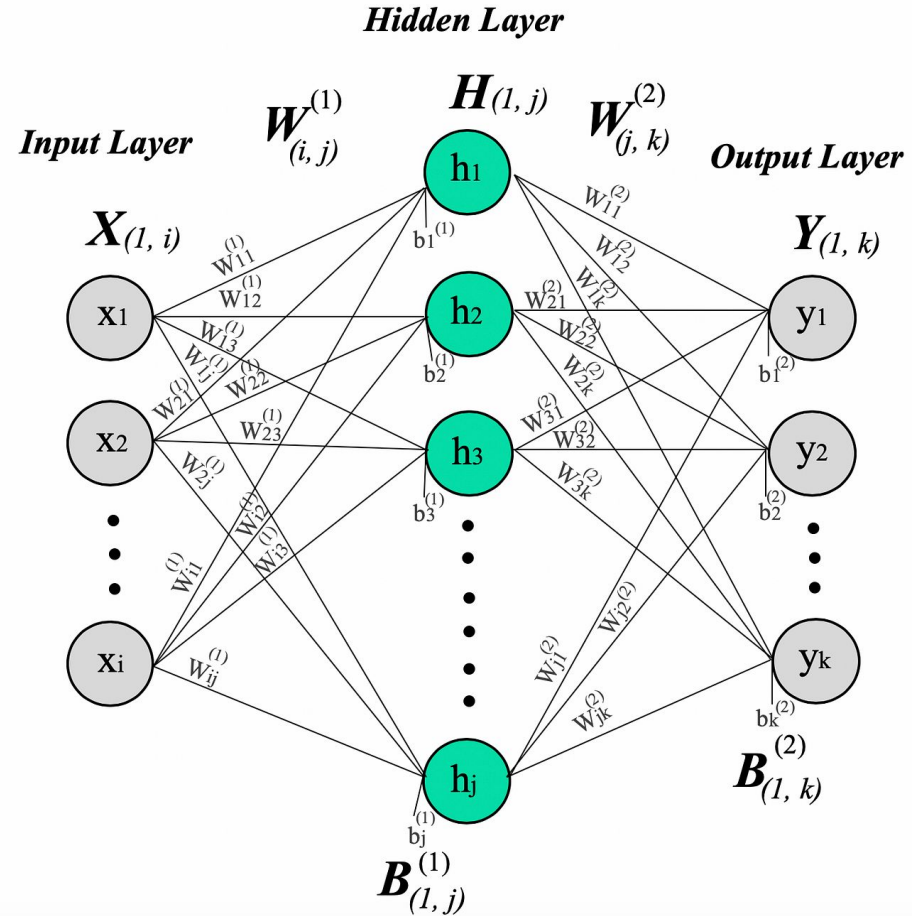


1 Neuron

- Cannot learn non-linear interactions between features.
- Cannot learn complex non-linearities between features and targets.
- This can be solved by using a hidden layer with more neurons.

NN – 3 Layers

- Input
- Hidden
- Output






NN – 3 Layer Lengths

- Input: same length as number of features.
- Hidden: any length we want (see later recommendations).
- Output:
 - 1 (typically linear) neuron for regression problems.
 - N neurons for classification problems with N classes.
 - Softmax if single-label.
 - Sigmoid if multi-label.
 - For binary classification, 1 sigmoid neuron suffices.


NN – 4 Layers

- Input
- Hidden
- Hidden
- Output






Deep NN (DNN)



NN – Common Activation Functions


- Linear
- Sigmoid
- Tanh
- ReLU
- Leaky ReLU
- ELU
- Softmax

Hidden
layers



NN – Common Activation Functions

- Linear
 - Sigmoid
 - Tanh
 - ReLU
 - Leaky ReLU
 - ELU
 - Softmax
- Output layer
- Output layer





NN – Loss

- Regression:
 - MSE
 - MAE
 - Huber loss (robust to outliers).




NN – Loss

- Regression:
 - MSE
 - MAE
 - Huber loss (robust to outliers).
- Classification:
 - Binary cross-entropy (same as log loss).
 - Sparse categorical cross-entropy (integer targets).
 - Categorical cross-entropy (one-hot encoded targets).

NN – Optimizing Loss

- Non-convex.
- Risk of getting stuck.
- Good learning rate.
 - Optimizers:
 - Adam
 - RMSprop
 - AdaGrad
- Good mini-batch for GD.



**STOCHASTIC
GRADIENT
DESCENT**



**SGD WITH
MOMENTUM**



RMSPROP



ADAM





NN – Epochs

- One epoch consists in one pass through all training samples.
- When separated in mini-batches, one epoch iterates over all batches.
- NNs require multiple epochs for a good optimization of the weights.
- Too many epochs can lead to overfitting.
- Too few epochs may result in underfitting.
- Early stopping can be applied.



Building NNs – Recommendations

- Number of layers:
 - One or two hidden for starting.
 - Maybe deeper for more complex data or task.



Building NNs – Recommendations

- Number of layers:
 - One or two hidden for starting.
 - Maybe deeper for more complex data or task.
- Neurons per layer:
 - It is common to start with a length between the input and the output.
 - Or even with more neurons than the input on the first layer.
 - Reducing the hidden sizes as the network goes deeper.



Building NNs – Recommendations

- Number of layers:
 - One or two hidden for starting.
 - Maybe deeper for more complex data or task.
- Neurons per layer:
 - It is common to start with a length between the input and the output.
 - Or even with more neurons than the input on the first layer.
 - Reducing the hidden sizes as the network goes deeper.
- Adam is normally a good start.




Building NNs – Number of Neurons

- More neurons can represent more complex patterns.
- More neurons are slower to train.
- Too many neurons can lead to overfitting.
- Too few neurons can result in underfitting.

NN – Example


- 4 Layers:
 - Input
 - Hidden 16 neurons
 - Hidden 8 neurons
 - Output 2 neurons with softmax
- Optimizer: Adam(0.01)
- Loss: categorical cross-entropy



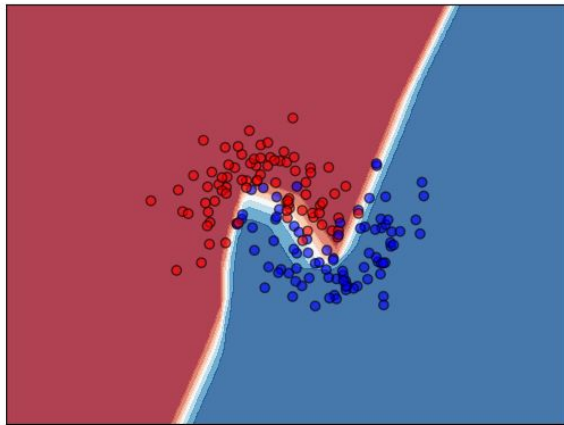


NN – Example


50 epochs



200 epochs



2,000 epochs





NNs – Libraries

- `sklearn.neural_network.MLPClassifier` or `MLPRegressor`
- Tensorflow
- Pytorch



Summary of Concepts

- Neural Network (NN).
- Deep NN (DNN).
- Weights (similar as LR coefficients).
- Activation functions.
- Input, Hidden, Output Layers.
- Forward/Backward propagation.
- Learning rate optimizers: Adam, RMSprop, AdaGrad.
- Epoch.