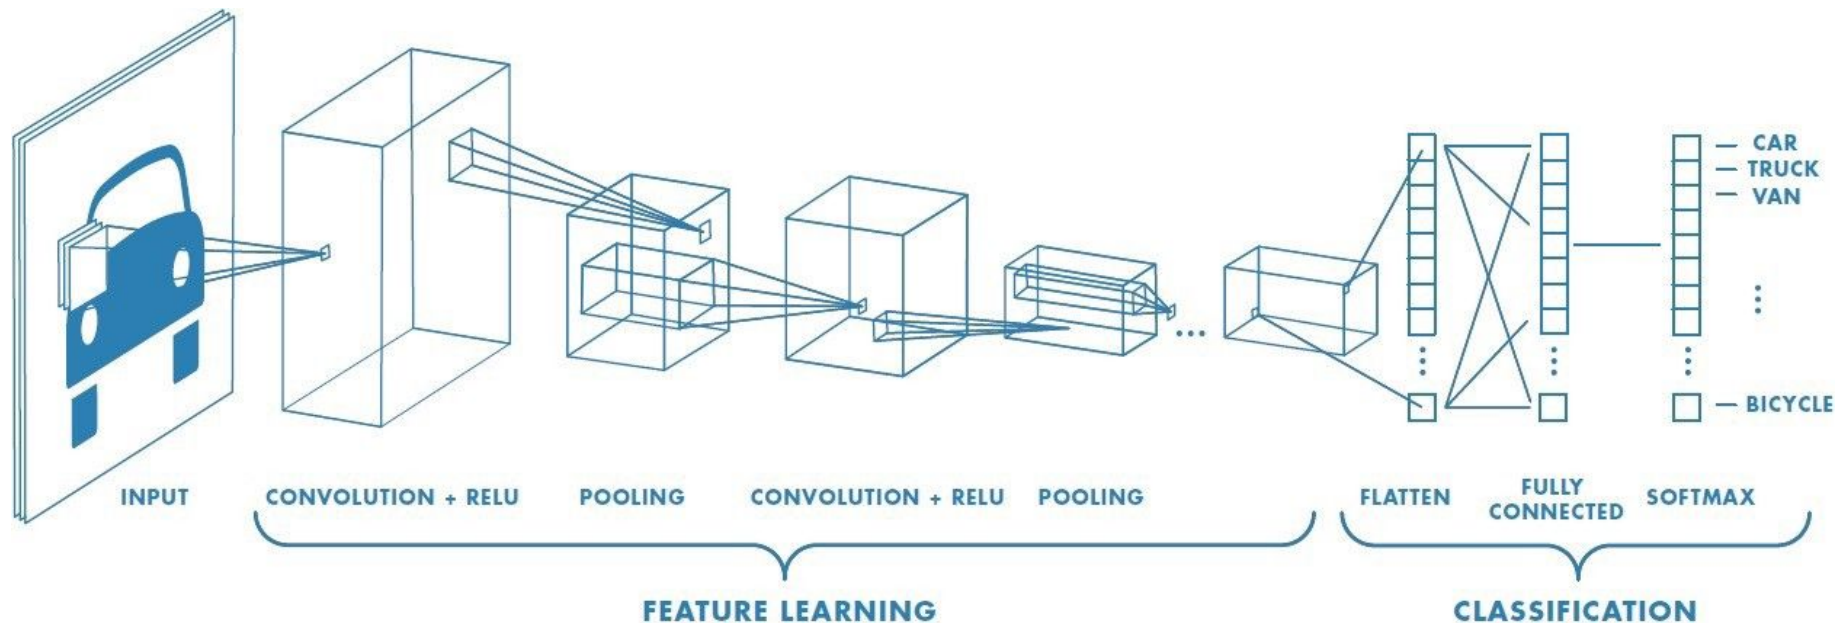
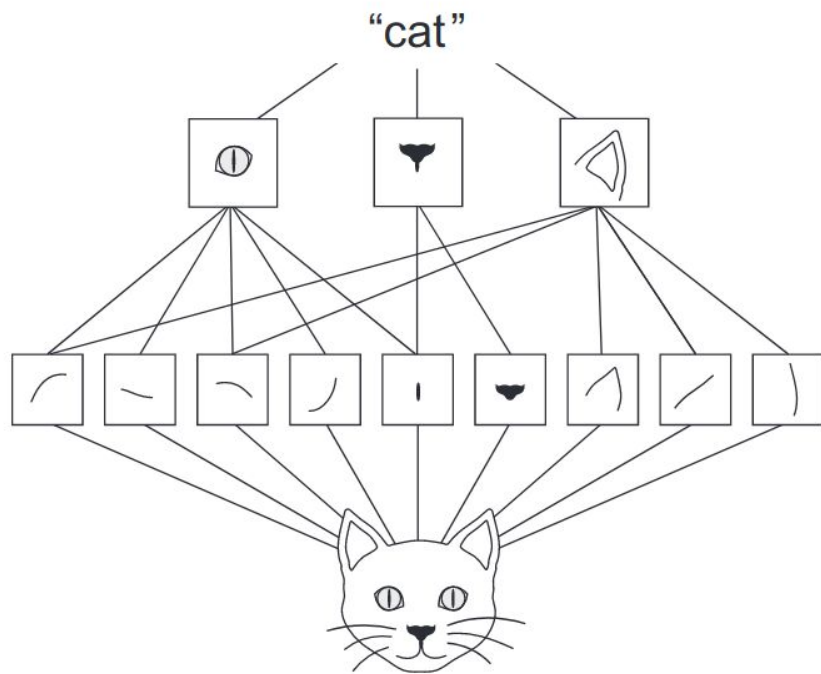





Convolutional Neural Networks (CNNs)

Convolutional Neural Networks (CNNs)






Convolution



Fuente: Bengio, Goodfellow y Courville (2016)

Convolution



0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Convolution



Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

Convolution


Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

ReLU



Output (3x3)

1		

Convolution


Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

ReLU



Output (3x3)

1	1	

Convolution


Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

ReLU



Output (3x3)

1	1	0

Convolution


Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

ReLU



Output (3x3)

1	1	0
1		

Convolution


Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

ReLU



Output (3x3)

1	1	0
1	3	

Convolution


Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

ReLU



Output (3x3)

1	1	0
1	3	1

Convolution


Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

ReLU



Output (3x3)

1	1	0
1	3	1
0		

Convolution


Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

ReLU



Output (3x3)

1	1	0
1	3	1
0	2	

Convolution




Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

ReLU



Output (3x3)

1	1	0
1	3	1
0	2	1

Convolution



Input (4x4)

0	0	0	0
0	1	1	0
0	0	1	0
0	0	1	0

Kernel (2x2)

1	1
0	1

Output (3x3)

1	1	0
1	3	1
0	2	1

Max-Pooling (2x2)




Input (3x3)


1	1	0
1	3	1
0	2	1

Output (2x2)


Max-Pooling (2x2)




Max-Pooling (2x2)



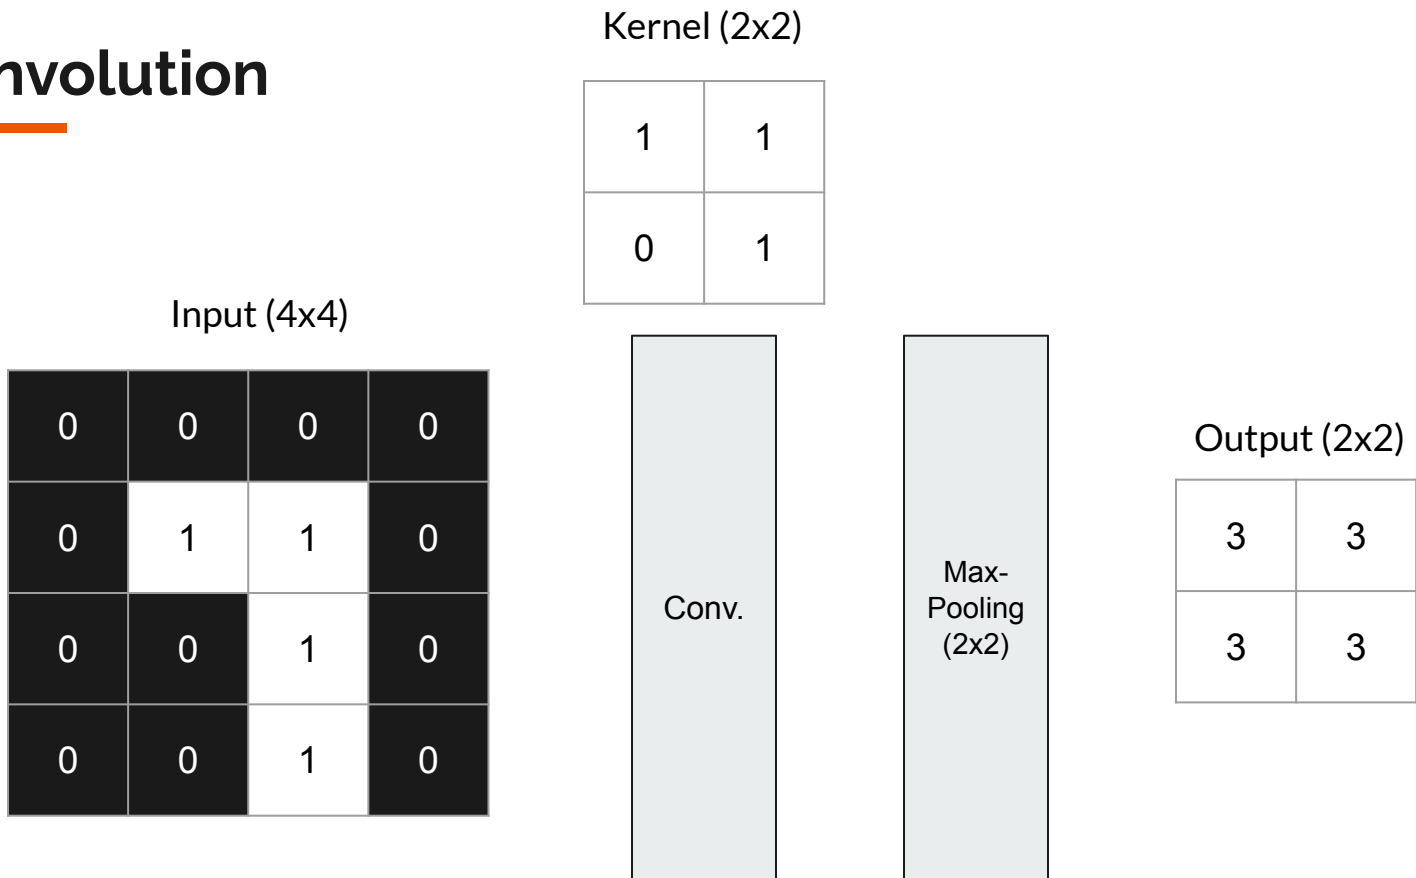
Max-Pooling (2x2)



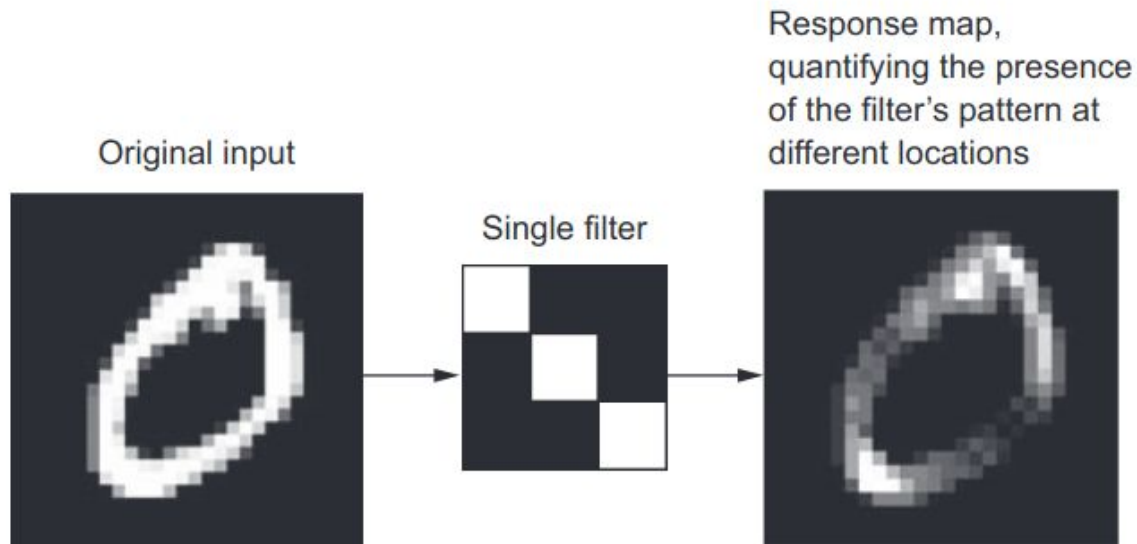
Max-Pooling (2x2)



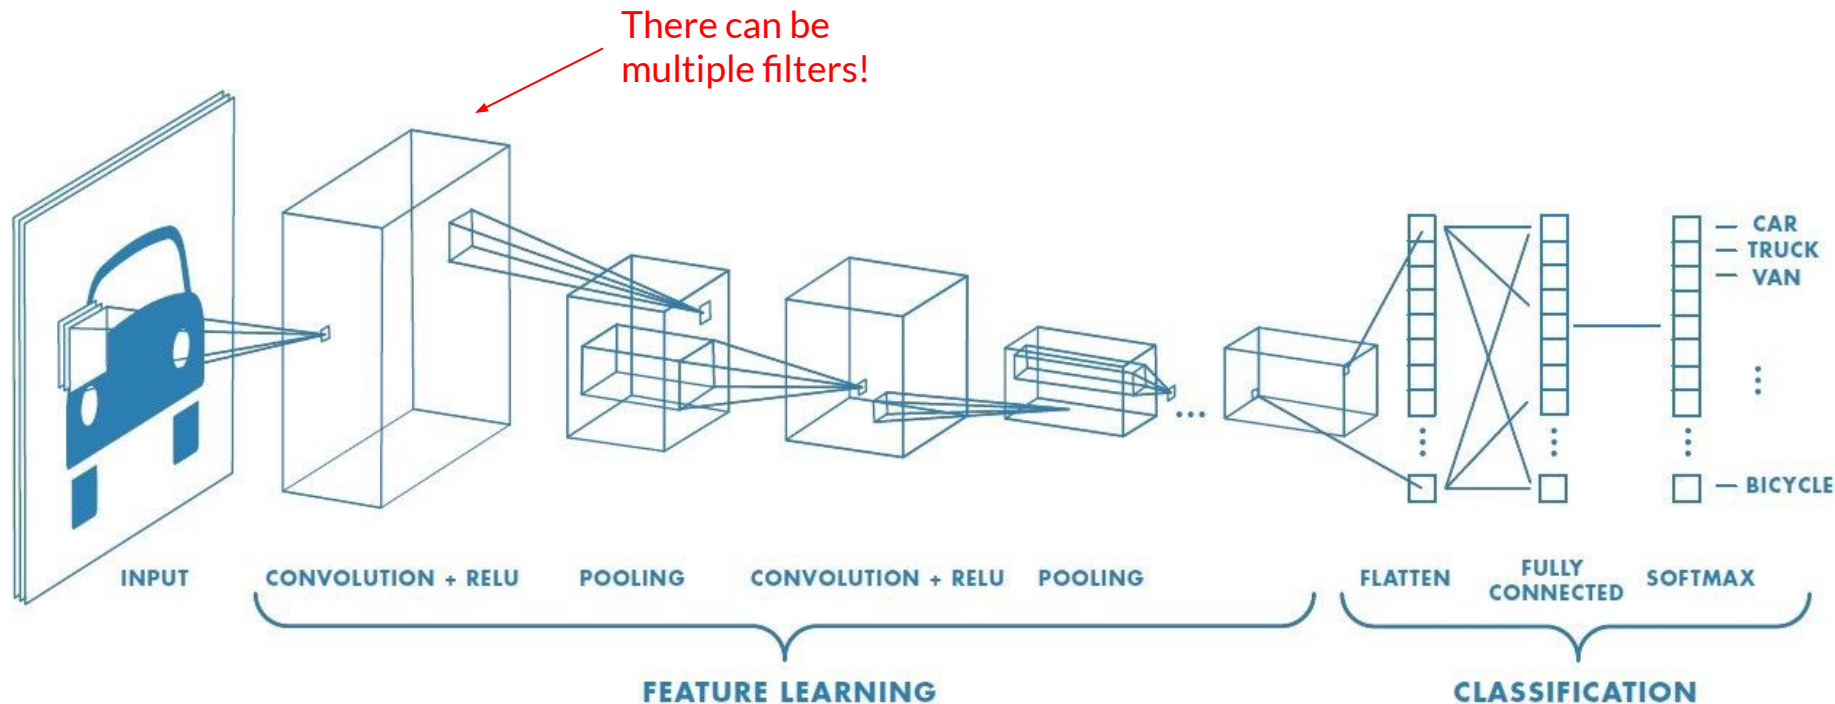
Convolution



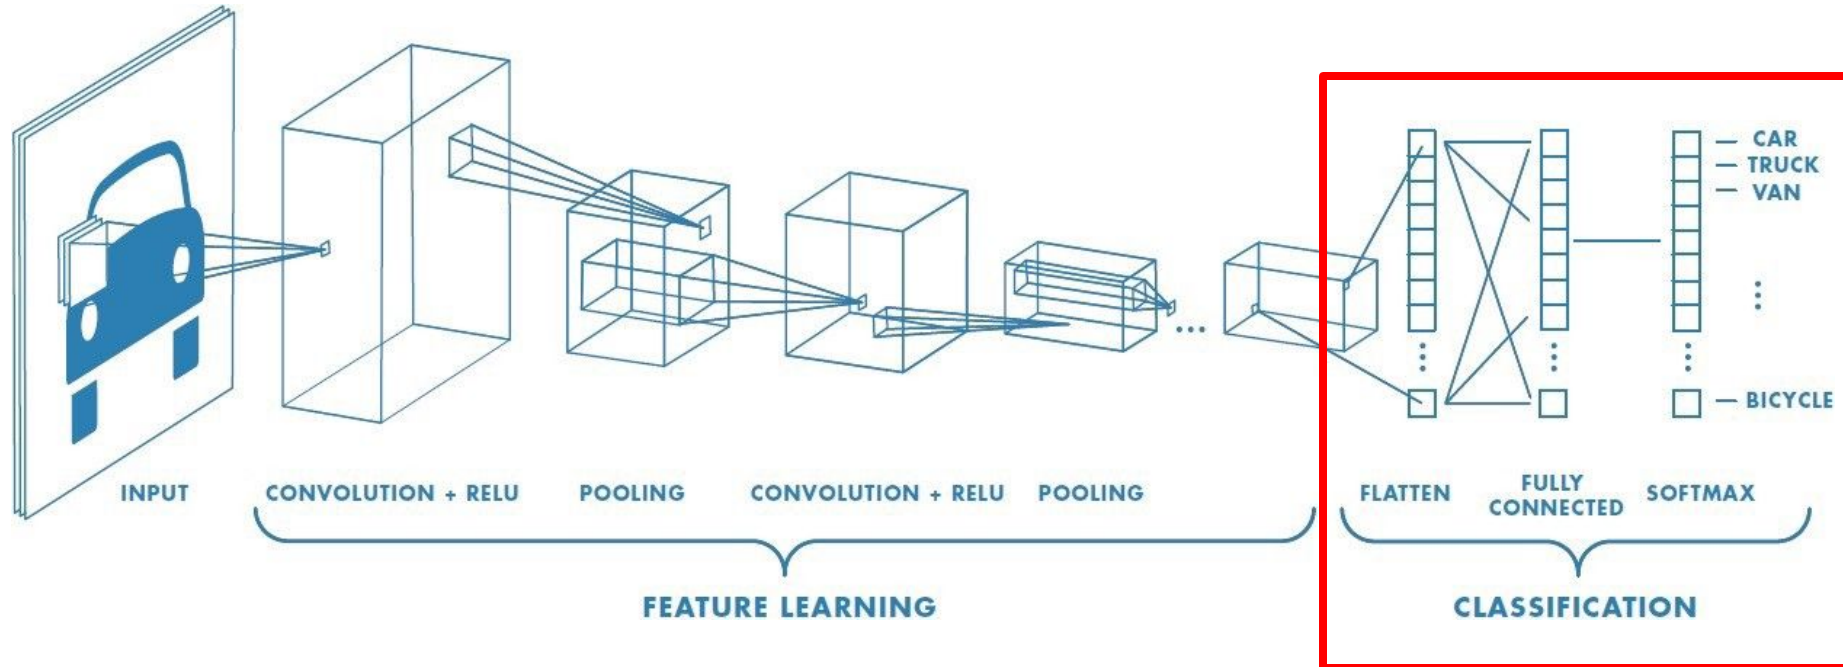
Convolution



Convolutional Neural Networks (CNNs)



Convolutional Neural Networks (CNNs)





Convolution

https://adamharley.com/nn_vis/cnn/2d.html

Padding


0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

Kernel

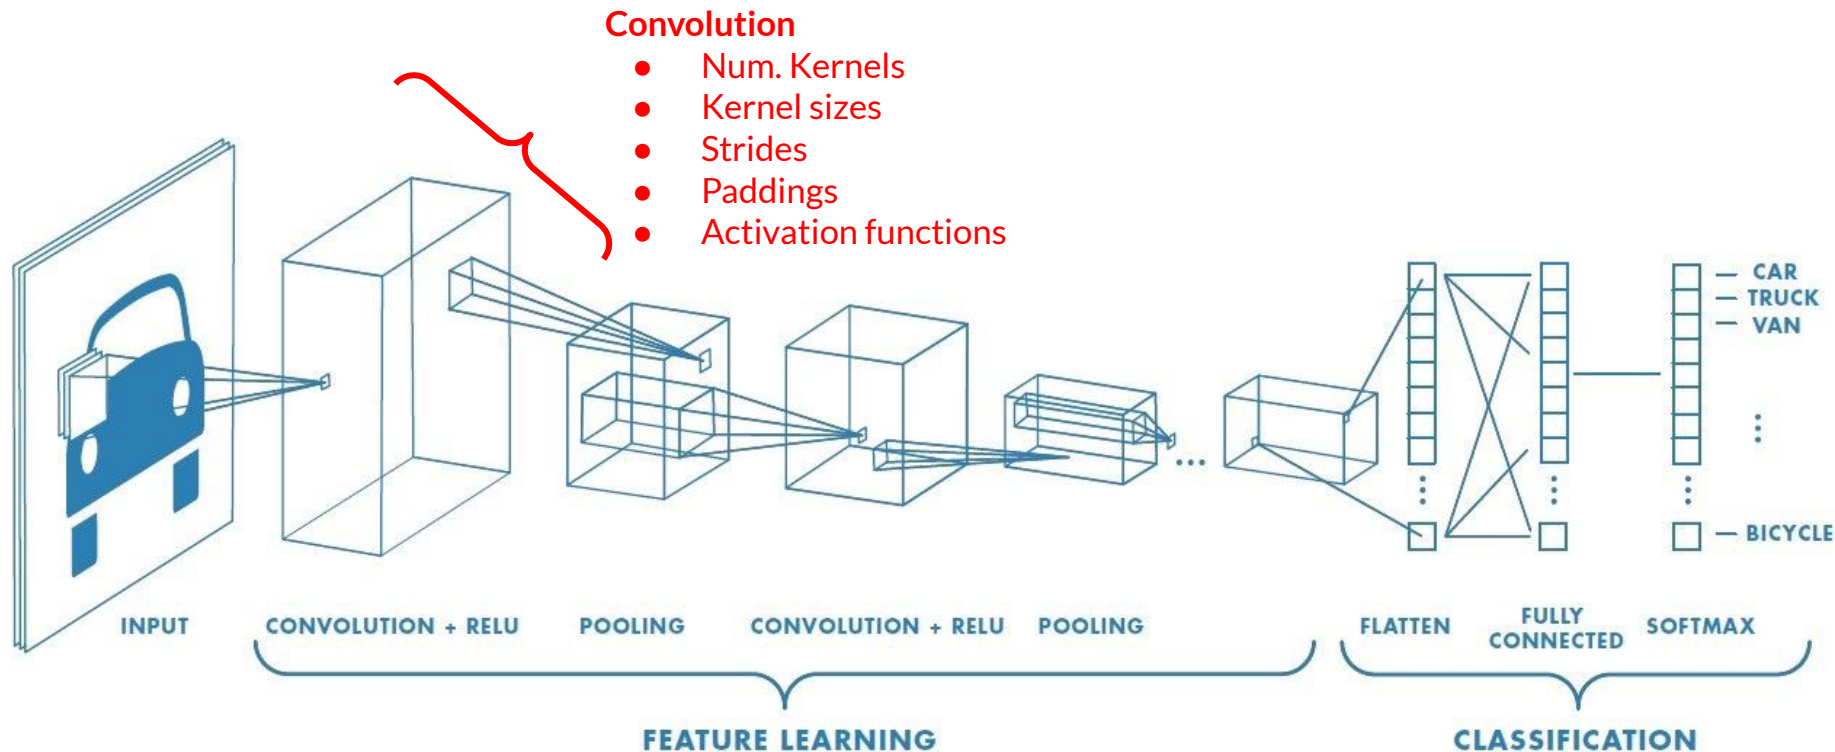
0	-1	0
-1	5	-1
0	-1	0

114				

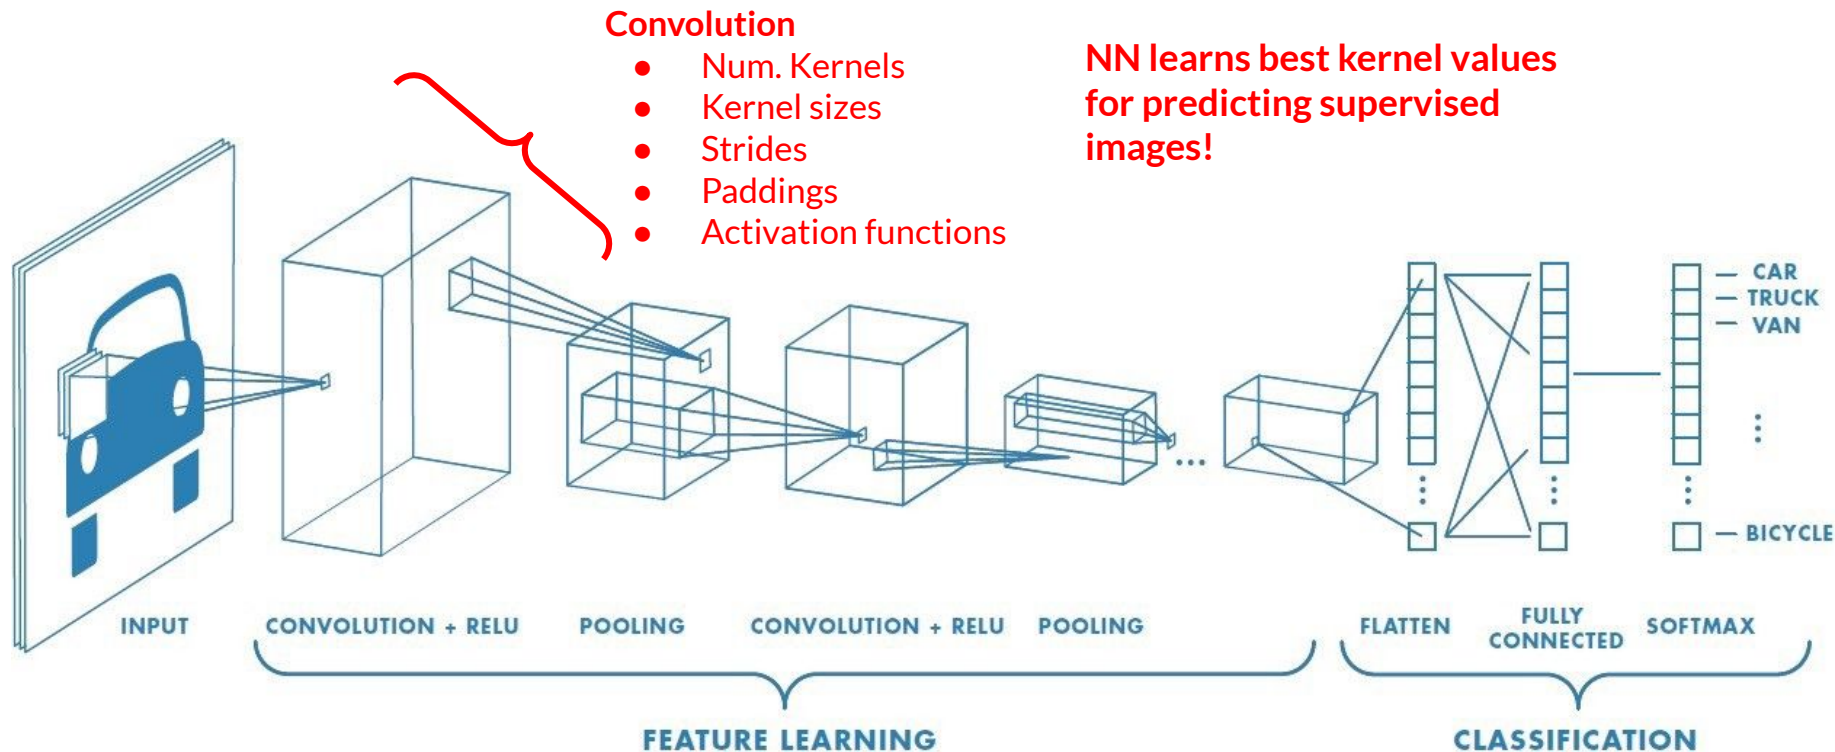
Stride = 2



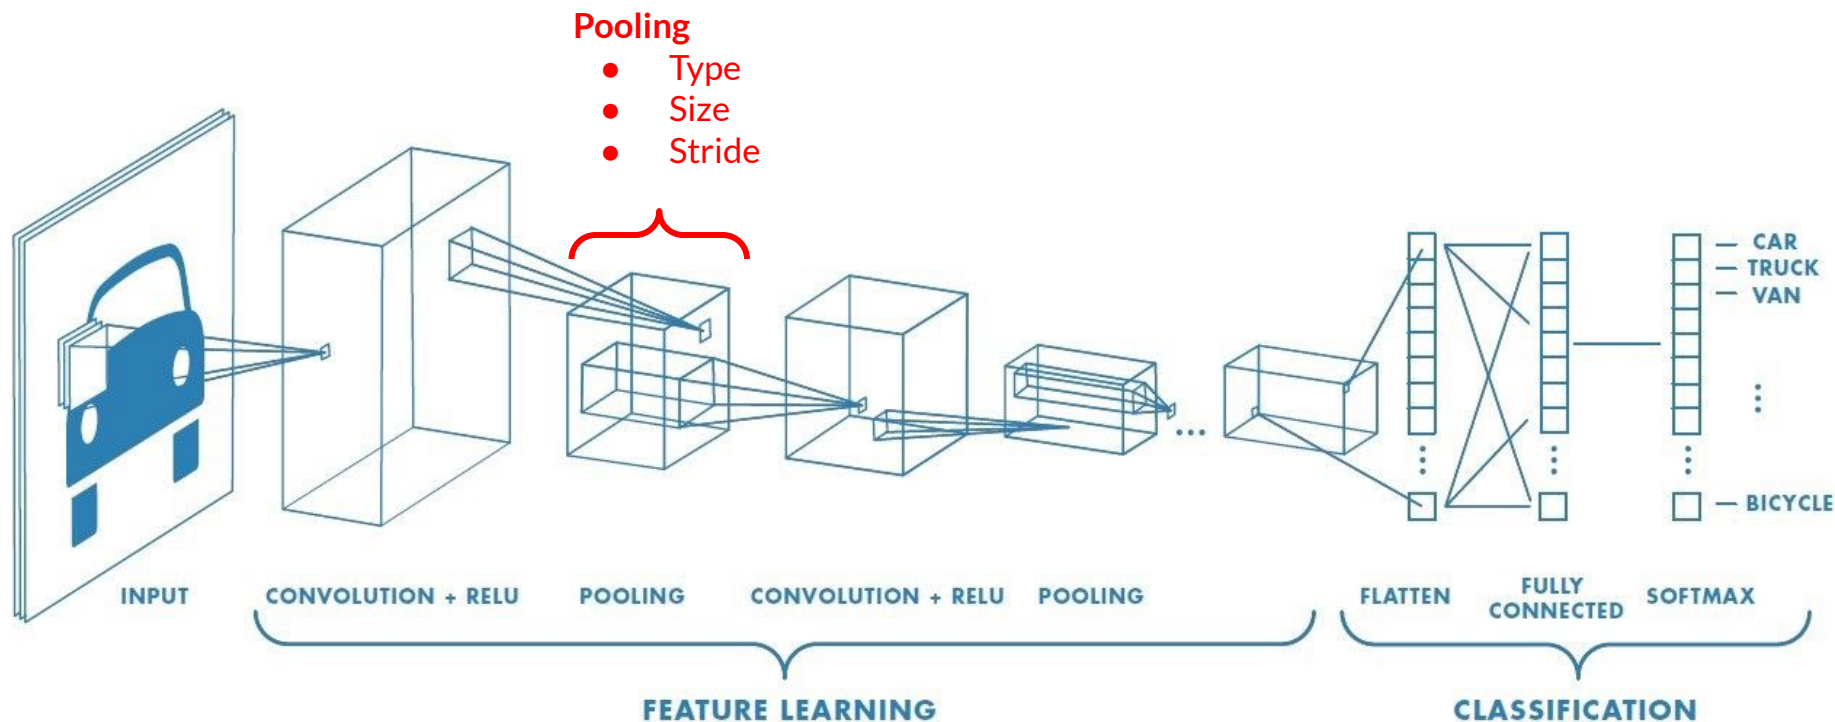
Convolutional Neural Networks (CNNs)



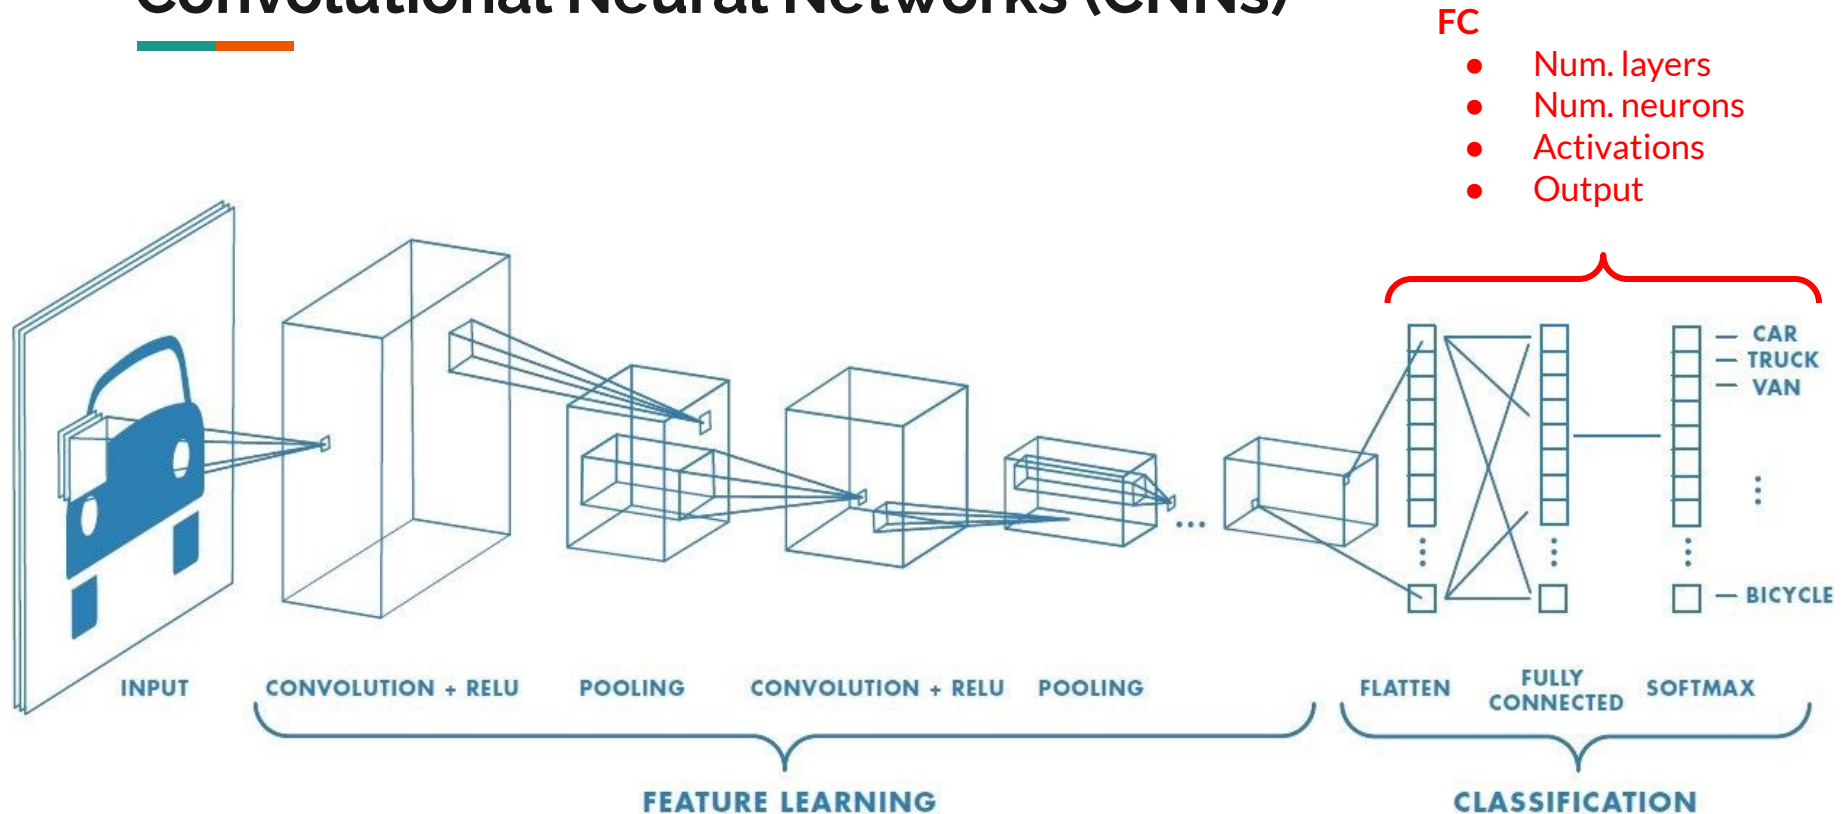
Convolutional Neural Networks (CNNs)




Convolutional Neural Networks (CNNs)




Convolutional Neural Networks (CNNs)




RGB Convolution



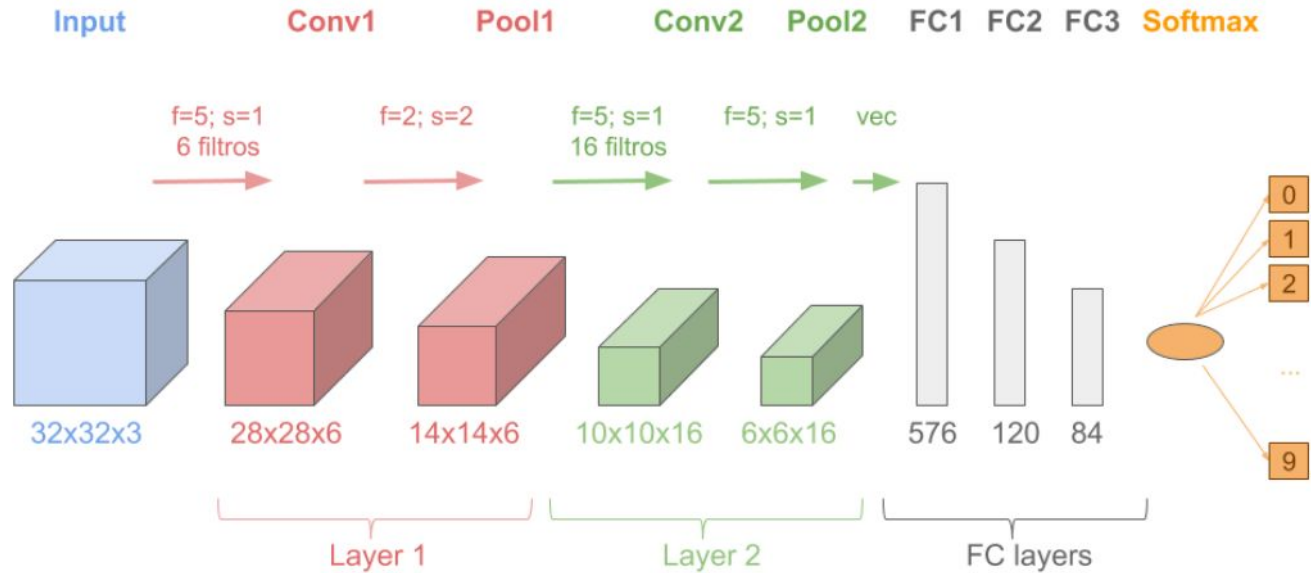
RGB Convolution



RGB Convolution



Inspired in LeNet-5



Fuente: LeCun y col. (1989)

Famous CNN architectures



Model name	Number of parameters [Millions]	ImageNet Top 1 Accuracy	Year
AlexNet	60 M	63.3 %	2012
Inception V1	5 M	69.8 %	2014
VGG 16	138 M	74.4 %	2014
VGG 19	144 M	74.5 %	2014
Inception V2	11.2 M	74.8 %	2015
ResNet-50	26 M	77.15 %	2015
ResNet-152	60 M	78.57 %	2015
Inception V3	27 M	78.8 %	2015
DenseNet-121	8 M	74.98 %	2016
DenseNet-264	22M	77.85 %	2016
BiT-L (ResNet)	928 M	87.54 %	2019
NoisyStudent EfficientNet-L2	480 M	88.4 %	2020
Meta Pseudo Labels	480 M	90.2 %	2021

Famous CNN architectures





<https://theaisummer.com/cnn-architectures/>

<https://keras.io/api/applications/>

224 x 224


INPUT



OUTPUT

VGG-16

VGG-16





Building CNNs – Recommendations

- Number of kernels
 - Small number in the first layers and increase it as you go deeper.
 - Early layers capture simple patterns and then you combine these in deeper layers.
 - Starting with 32 or 64 filters is common.



Building CNNs – Recommendations

- Kernel sizes
 - 3x3 and 5x5 are common options.
 - Smaller are usually preferred for capturing more specific details.
 - Sometimes, large kernels (e.g. 7x7) in the first layers for capturing general patterns.



Building CNNs – Recommendations


- Pooling layers
 - Generally used for reducing dimensions.
 - Using too many can imply losing valuable information.
 - 2x2 is very common, with greater poolings you risk losing valuable information.



Building CNNs – Recommendations

- Stride
 - A stride of 1 is common in convolutional layers.
 - A stride of 2 is common in pooling layers for reducing the dimensions of the image.

Interpretability



Misleading classification



“panda”

57.7% confidence


Misleading classification



“panda”

57.7% confidence

+ .007 ×



=




“gibbon”

99.3% confidence

Misleading classification



+ .007 ×



=



“panda”
57.7% confidence

“gibbon”
99.3% confidence



Resources

- <https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-convolutional-neural-networks>