



NNs More Concepts



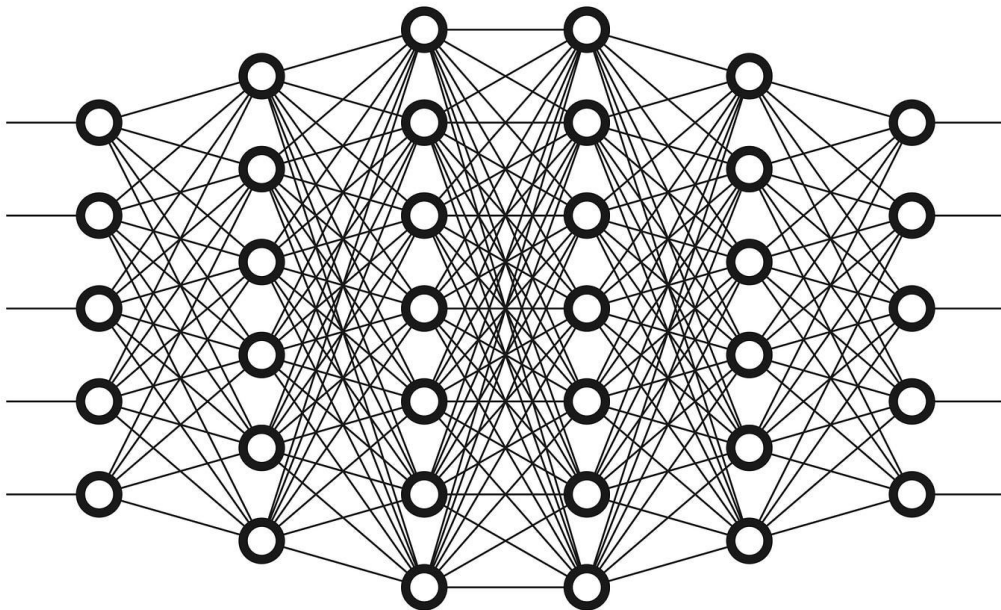
Common gradient problems

- Vanishing gradients.
- Exploding gradients.

Vanishing gradients



Significantly
smaller gradients





Vanishing gradients

- Symptoms:
 - Learning takes very long or stops completely.
 - Weights closer to the inputs do not change much.
 - Weights shrink exponentially.



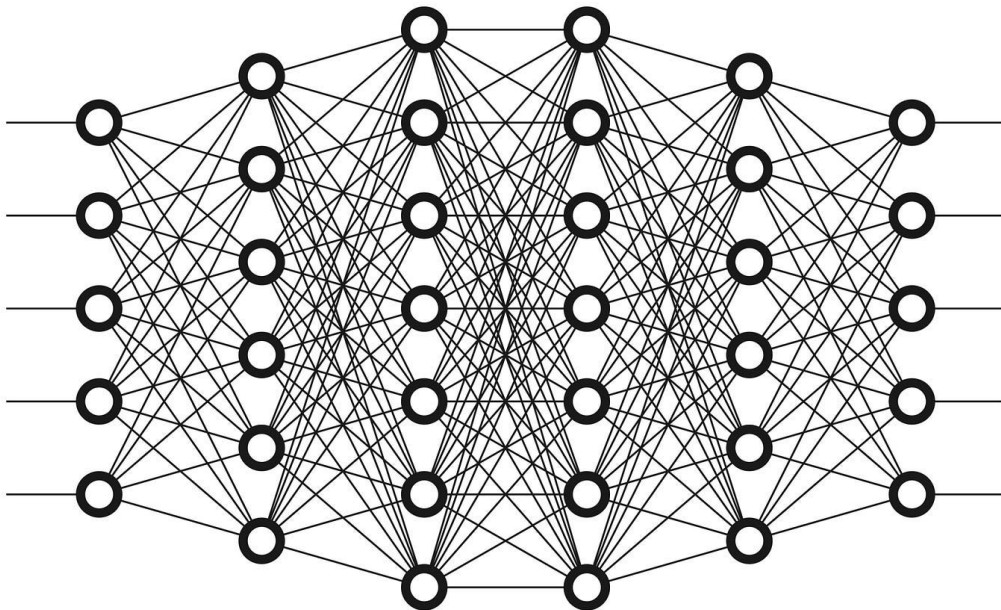
Vanishing gradients

- Symptoms:
 - Learning takes very long or stops completely.
 - Weights closer to the inputs do not change much.
 - Weights shrink exponentially.
- Solutions:
 - ReLU or similar.
 - Batch normalization.
 - Gradient clipping.
 - Reduce number of layers.

Exploding gradients



Significantly
larger gradients





Exploding gradients

- Symptoms:
 - Exponential weight growth.
 - Weights overflow (NaN).
 - Weights closer to the inputs are larger.



Exploding gradients

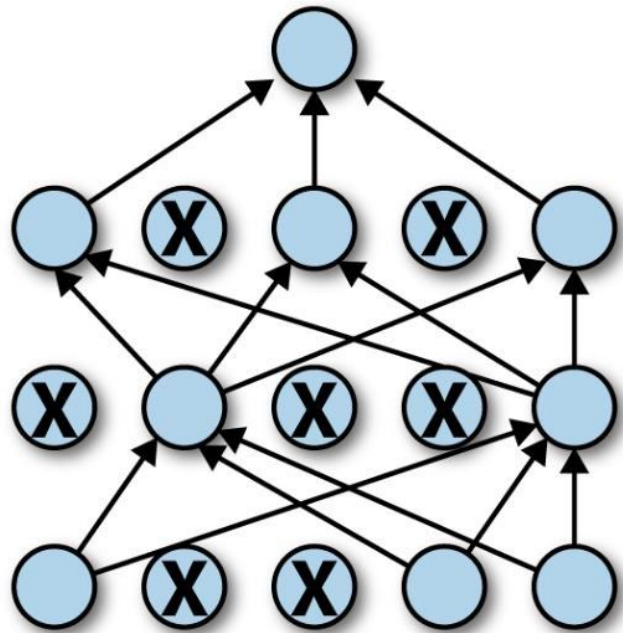
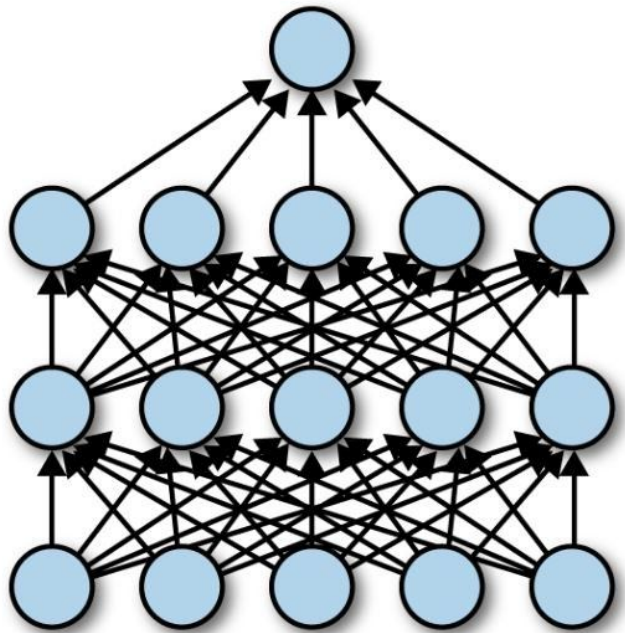
- Symptoms:
 - Exponential weight growth.
 - Weights overflow (NaN).
 - Weights closer to the inputs are larger.
- Solutions:
 - Proper weight initialization.
 - Gradient clipping.
 - Reduce the number of layers.



Regularization (for avoiding overfitting)

- Early stopping.
- Weight initialization.
- Dropout.
- Batch normalization.

Dropout





Dropout layer – Keras

```
nn_model = Sequential([
    Input(X_train.shape[1]),
    Dense(128, activation='relu'),
    Dropout(0.2),
    Dense(64, activation='relu'),
    Dropout(0.2),
    Dense(32, activation='relu'),
    Dense(len(class_names), activation='softmax')
])
```

- Randomly set 20% of parameters to 0 at each mini-batch.
- Applied to the previous Dense layers.



Dropout layer – Keras

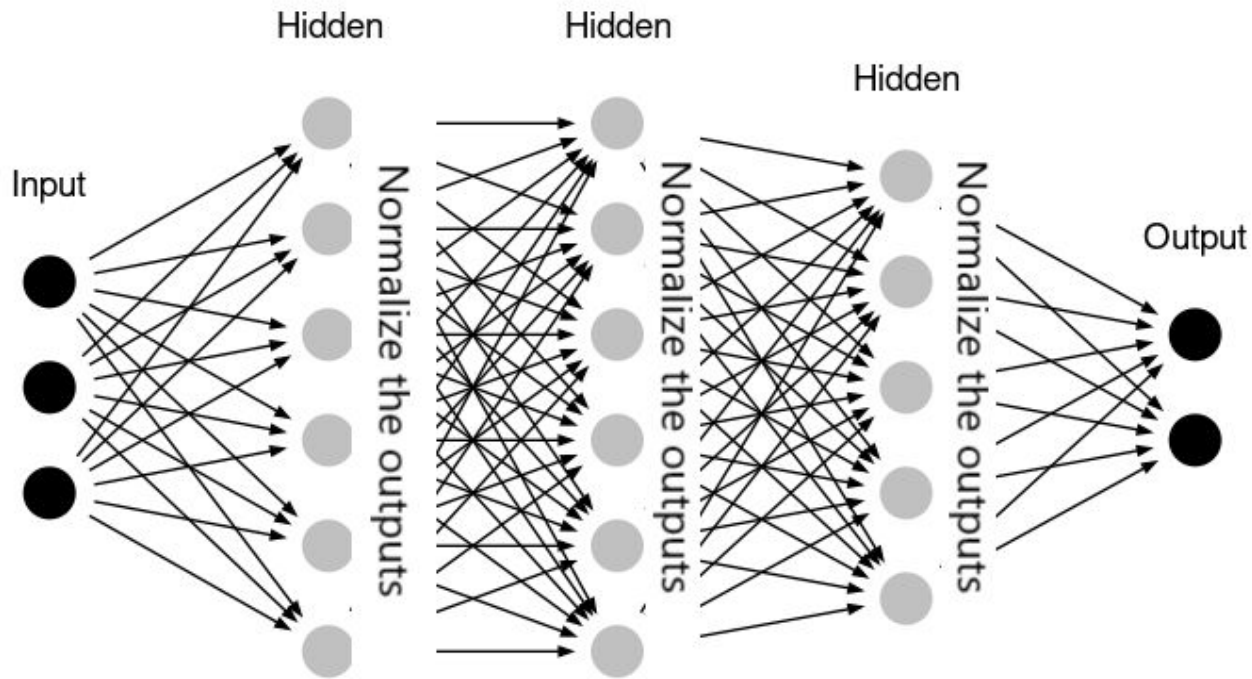
```
cnn_model = Sequential([
    Input(shape=(256, 256, 1)),
    Rescaling(1./255),
    Conv2D(filters=16, kernel_size=2, strides=1, activation='relu'),
    Dropout(0.1),
    MaxPooling2D(pool_size=2, strides=None),
    Conv2D(filters=32, kernel_size=2, strides=1, activation='relu'),
    Dropout(0.1),
    MaxPooling2D(pool_size=2, strides=None),
    Flatten(),
    Dense(32, activation='relu'),
    Dense(16, activation='relu'),
    Dense(len(class_names), activation='softmax')
])
```



Dropout layer – Recommendations

- Dropout rate:
 - Simple NNs: 20% or 30%
 - Complex NNs: 40% or 50%
 - CNNs: 10% or 20%
- Not necessarily used in each layer.

Batch normalization





BatchNormalization layer – Recommendations


- Don't use it when:
 - Small mini-batches.
 - RNNs.

Data augmentation

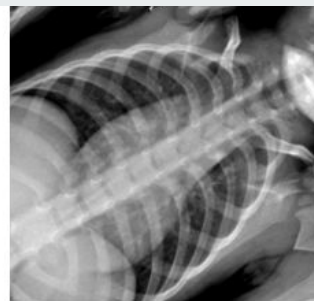
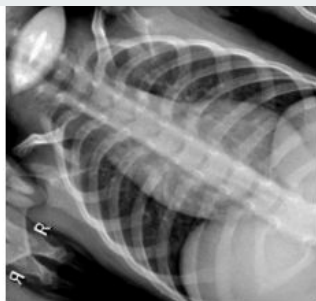


**Image
Augmentation**

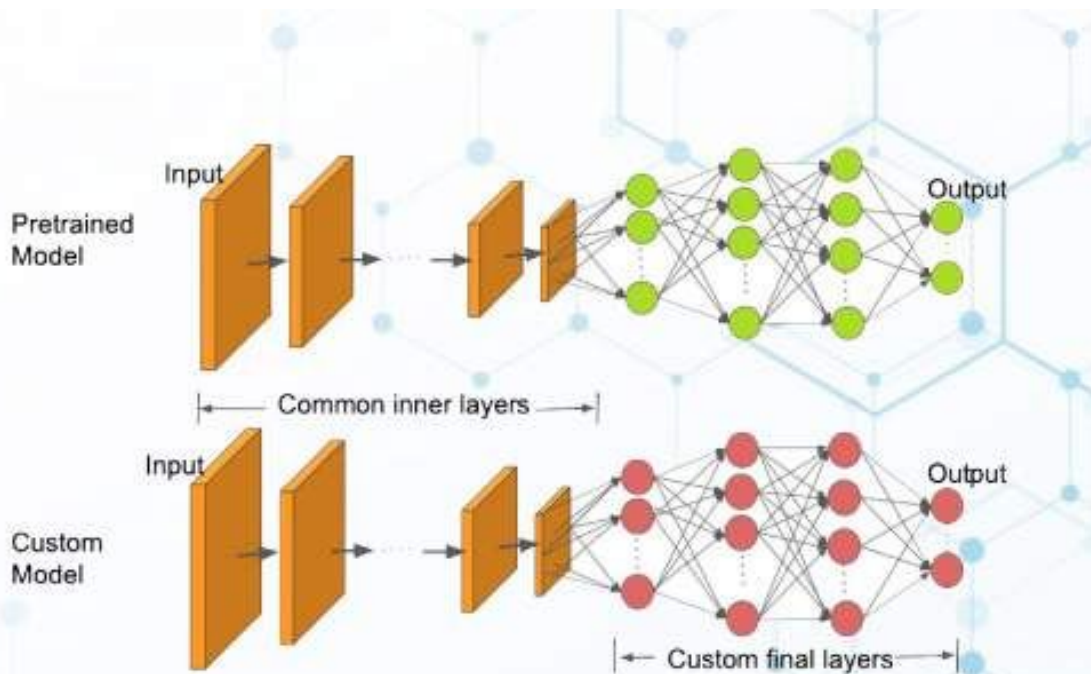




```
cnn_model = Sequential([
    Input(shape=(256, 256, 1)),
    Rescaling(1./255),
    layers.RandomFlip("horizontal_and_vertical"),
    layers.RandomRotation(0.2),
    Conv2D(filters=16, kernel_size=2, strides=1, activation='relu'),
    MaxPooling2D(pool_size=2, strides=None),
    Flatten(),
    Dense(32, activation='relu'),
    Dense(len(class_names), activation='softmax')
])
```



Transfer Learning



```
from tensorflow.keras.applications import MobileNetV2
```

```
base_model = MobileNetV2(  
    weights='imagenet',  
    include_top=False,  
    # Not the default size, but we can use it anyway  
    input_shape=(256, 256, 3))
```

```
for layer in base_model.layers:  
    layer.trainable = False
```

```
# Create a new sequential model
cnn_model = Sequential()

# # Resize the images to 224x224
# cnn_model.add(Resizing(224, 224))

cnn_model.add(Input(shape=(256, 256, 1)))

# Rescale the grayscale images
cnn_model.add(Rescaling(1./255))

# Expand the 1 channel input to 3 channels
# to match MobileNetV2 input requirements
cnn_model.add(Conv2D(3, (3, 3)))

# Add all the MobileNetV2 layers
cnn_model.add(base_model)

# Add your custom layers
cnn_model.add(Flatten())
cnn_model.add(Dense(32, activation='relu'))
cnn_model.add(Dense(len(class_names), activation='softmax'))

# Compile the model
cnn_model.compile(optimizer=Adam(learning_rate=0.001),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])
```